



Qognito.io

OCT
2026

Human oversight,
costs, time actually
reallocated: the
blind spots in your
AI projects' ROI.

AI'S PHANTOM BILL

**THE MIRAGE OF TOTAL DELEGATION—AND THE
SAGA METHOD FOR STOPPING IT FROM
DRAINING YOUR CASH.**



The speed of
generative AI
masks an
explosion in your
oversight
workload.

PHASE 1: THE DIAGNOSIS

THE MIRAGE OF TOTAL DELEGATION

1. The boardroom syndrome

You have all seen it: the perfectly choreographed demonstration.

In seconds, in front of the executive committee, an artificial intelligence tool analyses a 50-page report, extracts the strategic issues and drafts a summary ready for distribution. The execution is lightning-fast. The result is impressive. The conclusion seems obvious: a process that used to take your teams hours has been reduced to a single click. The promise of a huge return on investment (ROI) is on the table.

Yet a few months later, the reality of deployment in production tells a very different story. Productivity gains fail to translate into cash. Teams complain of a heavier workload. Proofs of concept (POCs) drag on without ever scaling up.

Que s'est-il passé entre la démonstration éblouissante et l'enlisement opérationnel ?

What happened between the dazzling demonstration and the operational quagmire? The business was blinded by the mirage of total delegation. It confused AI's ability to generate text with its ability to take real responsibility for business work.



2 . Human oversight and inference debt: two costs to distinguish

The first analytical mistake is to assume that the time the machine saves on drafting makes the workload disappear.

In reality, generative AI does not eliminate cognitive work: **it shifts and transforms it.**

Generating text has become free and instantaneous. But checking its accuracy, rooting out hallucinations—those moments when AI invents a fact with deceptive confidence—and making sure the business perspective is respected takes human effort that cannot be eliminated.

This work constitutes human oversight and rework effort. If an AI system produces 1,000 monitoring summaries a month, but 20% contain inaccuracies that an expert must manually correct after the fact, the initial time saving is wiped out by the rework.

We reserve the term **inference debt** for the avoidable recurring cost of having the model rediscover known rules that could have been captured for reuse in its execution environment. Necessary human oversight is therefore not, in itself, inference debt. A decision-maker may believe they have automated a task when they have turned their senior analysts into an algorithm's proofreaders and fact-checkers.



3 . Stop paying for effort. Demand results.

To regain financial and strategic control of your automation projects, you must first redefine **the unit by which success is measured**.

Today, many service providers and IT departments present misleading performance indicators that charge for the machine's effort, rather than the value delivered to the business. If an AI needs three attempts—and consumes three times as much computing power—to churn out a summary that a human will still have to edit heavily, the overall cost balloons without adding a single useful unit to the tally.

The only valid criterion for an executive is an **accepted, actionable deliverable**.

- Information has value only if it informs or triggers a decision.
- The basic unit of your ROI is information accepted by a human against your strict quality criteria—not the document processed.

As long as you assess your AI projects by generation speed rather than the full cost of validation, your initiatives will remain budget sinkholes. Ask yourself: in your last POC, who timed the hours your experts spent reviewing the machine's output?

The challenge is therefore no longer technological: it is a matter of architecture and governance. It means structuring what we call the operational harness, the only safeguard that can ensure AI does not generate more debt than value.



Human oversight
accounts for
approximately
70% of the total
cost.

PHASE 2 : THE COST MODEL

THE PHANTOM BILL AND THE REALLOCATION IMPERATIVE

1 . The illusion of the one-sided bill (C1 vs TCO)

An enthusiastic project manager presents an investment proposal to the executive committee. The monthly API bill—access to the AI model—comes to €600. Against this, they set the historical cost of manual processing: €36,000. The subtraction is spectacular; deployment seems a mathematical no-brainer.

Yet this dashboard is a dangerous fiction because it compares fundamentally different things. Comparing the raw cost of inference with a full payroll bill is like comparing the price of jet fuel with the total cost of an airline fleet. Your €600 API bill buys the fuel. But it does not buy the aircraft—the software infrastructure—the control tower—governance—or the pilot—the human who corrects the machine's deviations.

To move beyond magical thinking, an executive must demand a **Total Cost of Ownership (TCO)** assessment that rigorously separates costs into three distinct taxes:

A. [C1] The Cognitive Tax : This is the price of APIs: the charge for each request sent to the language model. Remember that this tax applies to successes and errors alike, including correction loops.



B. [C2] The Infrastructure Tax : The API does not run on its own. Orchestration platforms, vector databases, hosting servers and cybersecurity integration all have to be paid for.

C. [C3] The Human Oversight Tax : This is the dark mass of the iceberg below the waterline. It includes document review time, manual correction of hallucinations after the fact, and continuous auditing of the agent's rules. In our SAGA simulation, **this tax accounts for approximately 70% of the total cost**; this ratio is not an observed average across enterprise deployments.

Illustration from the SAGA teaching case: 200 hours of residual human work, valued at €60/hour, amount to €12,000 out of a total cost of €16,500—approximately 73%. This ratio describes that educational case, not enterprise deployments as a whole.

The real challenge for a finance and administration department is not to list these taxes, but to know how to audit software vendors' proposals to uncover the true cost of C3.

This is one of the core skills taught in the SAGA-IA programme.



2 . Freed-up time without a reallocation plan erodes cash (The empty warehouse syndrome)

Imagine that, even after including the substantial human oversight tax (C3), the system frees up a net 400 hours a month for your teams. On paper, that is a win.

In financial reality, this is the **empty warehouse syndrome**. Saving time without a concrete plan to use it is like renting a warehouse with no goods to store in it. You pay the rent—the C1 and C2 taxes—but the empty space earns nothing.

The time freed up by AI is often fragmented: 10 minutes here, 5 minutes there. If employees use it to check their emails or polish secondary presentations, the business gains nothing. Worse, it has increased its spending on APIs and servers while its payroll has remained unchanged: **automation has worsened its immediate cash position**.

For those saved hours to have economic value, they must be assessed through the Reallocation Rate. This is the share of freed-up time that the business can demonstrate has been reused for a revenue-generating task—a new contract or billable work—or for concrete cost savings, such as cost avoidance or a hiring freeze.

From this measure, we calculate the **Alpha Threshold**, the proprietary profitability indicator you will learn to model in the SAGA-IA programme.



This is the equation that determines whether an AI project survives. The project is profitable only if the financial value of the reallocated time exceeds the cost of the technological iceberg ($C1 + C2 + \text{Amortisation}$).

In our SAGA educational simulation, the capacity balance turns positive when the value of the net time reallocated exceeds €4,500 in technical expenditure and amortisation: €600 for APIs + €2,400 for infrastructure + €1,500 in amortisation. Methodological reminder: residual human work has already been deducted when calculating the net time freed up, so it must not be subtracted a second time. Moreover, the value of reallocated capacity does not automatically constitute a gross cash saving.

3 . Profitability depends on industrial-scale volumes (Stall speed)

Finally, the profitability of this equation is dictated by volume. AI is not inherently profitable “from the first document”. The fixed costs of governance and infrastructure— $C2$ plus part of $C3$ —impose a stall speed: a minimum viable volume.

If your department aims to process a small number of operations, fixed costs will completely overwhelm the modest marginal gain from drafting a few texts. Deploying autonomous agents for minor administrative tasks makes no economic sense. For low-volume microtasks, conventional or scripted automation remains far more economical than an autonomous agent, at an equivalent level of output quality. An AI safety harness is substantial: **it requires industrial-scale volumes or very high-value tasks to pay for itself.**



Stop funding
technology
demonstrations.
Build viable
services.

PHASE 3 : REGAIN CONTROL

YOUR ARCHITECTURE AND GOVERNANCE AUDIT

1 . The absence of an operational harness turns your investments into chronic cost centres

If your artificial intelligence initiatives struggle to demonstrate profitability, it is not because the technology is immature. It is because your projects' architecture was designed to shine in the boardroom, rather than withstand real-world operational friction.

The absence of a rigorous operational harness, combined with an inability to enforce a meaningful Reallocation Rate, turns your technology investments into chronic cost centres. Inference debt builds up silently, masked by the illusion of execution speed.

The time for technological wonder is over; the time for management discipline has arrived. It is time to hold your AI projects to the same profitability standards as any other strategic investment.



2 . Safeguard your investments through a rigorous examination of TCO

Qognito provides the means to work with decision-makers—chief executives, CFOs and CIOs—to redesign or build architectures for delegating work to AI. My approach rests on a rigorous examination of Total Cost of Ownership (TCO) and residual human effort, rather than vendors' promises. A

As a technical expert and AI strategist, I provide the toolbox you need to regain control of your deployments. Qognito supports you through two forms of intervention:

A. SAGA-IA Training (AI Agent Architecture and Governance Strategist) :

The aim is to equip you with the intellectual shield you need to stop technology integrators dictating your ROI projections, and to block toxic projects before they consume your cash. SAGA-IA is the intensive programme designed to equip your leadership teams and project managers with the analytical frameworks presented in this document: C1/C2/C3 and the Alpha Threshold.

You will then be able to audit your own projects independently and make a reasoned decision about whether to continue or stop them.

(Progress update: the two sections covering diagnosis and the economic model have been written. The final syllabus and presentation materials are being completed.)



B. Technical Design of the Operational Harness :

For critical projects requiring direct involvement, close collaboration makes it possible to design the security and governance framework: service-level agreements (SLAs), data-flow mapping and approval gates. The objective is to provide a technical guarantee that AI will not publish, distribute or validate sensitive data without explicit, traceable human authorisation.

4 . Move from illusion to operation

Do not let an attractive API bill hide rising payroll costs. Automation must serve your cash position, not the other way round.

Take control of your AI delegation strategy today.

CONTACT:



Website
[Qognito.io](https://qognito.io)



Email
bguarisma@qognito.io



LinkedIn
[Qognito](https://www.linkedin.com/company/qognito)



LEGAL INFORMATION & REFERENCES

This strategic discussion paper draws on the research and framework behind SAGA-IA (Stratège en Architecture & Gouvernance des Agents IA — AI Agent Architecture and Governance Strategist), developed by Qognito. (SAGA-IA source materials: *Section 1 — Le Renversement : le diagnostic d'inférence*; *Section 2 — Le modèle économique de l'IA*. The numerical cases are educational simulations.)

QOGNITO (SASU) share capital €1,000.00 — Registered office: 151 Avenue Félix Faure, 92800 Puteaux — Registered with the Nanterre RCS under no. 130 175 730 — VAT no.: FR07130175730.

QOGNITO — Tous droits réservés © 2026